

## Deep Learning

### Assignment:

Apply deep learning technique using the h2o package to the Wisconsin Breast cancer data set [https://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+\(Original\)](https://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+(Original)) .

First, I prepared the environment, loaded required libraries, connected to the local h2o cluster, and loaded the data set using the following code:

```
> ### Deep Learning
> ###Wisconsin Breast Cancer Dataset
> ###Dataset source: https://archive.ics.uci.edu/ml/machine-learning-databases/breast-cancer-wisconsin/
> rm(list=ls()) #Clear the environment
> setwd("YOUR_PATH") #Set working directory for the assignment
> getwd() #Check working directory
[1] "YOUR_PATH"
>
> #Libraries
> library(h2o) #load library
> localH2O = h2o.init(nthreads = -1) #connect to a local cluster
Connection successful!
```

R is connected to the H2O cluster:

```
H2O cluster uptime:      6 hours 57 minutes
H2O cluster timezone:    America/Denver
H2O data parsing timezone: UTC
H2O cluster version:     3.20.0.8
H2O cluster version age:  2 months and 23 days
H2O cluster name:        Rodneyweakly
H2O cluster total nodes: 1
H2O cluster total memory: 3.20 GB
H2O cluster total cores: 8
H2O cluster allowed cores: 8
H2O cluster healthy:     TRUE
H2O Connection ip:       localhost
H2O Connection port:     54321
H2O Connection proxy:    NA
H2O Internal Security:   FALSE
H2O API Extensions:      Algos, AutoML, Core V3, Core V4
R Version:               R version 3.5.1 (2018-07-02)
```

```
> library(tidyverse)
> library(ggplot2)
> library(caret)
> #Load data
> bcancer <- read.table("breast-cancer-wisconsin.data", header = FALSE, sep = ",", stringsAsFactors = FALSE)
```

To verify that the data loaded correctly, I looked at the internal structure of the data frame, and the first few and the last few observations. The data frame contains 699 observations of 11 variables describing various cell measurements and the presence or absence of the breast cancer diagnosis for each observation. All variables loaded as integers, except for the variable #7 with loaded as character array suggesting that there some missing or inappropriately formatted observations there.

```
> #EDA and data pre-processing
> str(bcancer)
'data.frame':      699 obs. of  11 variables:
 $ v1 : int  1000025 1002945 1015425 1016277 1017023 1017122 1018099 1018561 1033078 1033078 ...
 $ v2 : int   5 5 3 6 4 8 1 2 2 4 ...
 $ v3 : int   1 4 1 8 1 10 1 1 1 2 ...
 $ v4 : int   1 4 1 8 1 10 1 2 1 1 ...
 $ v5 : int   1 5 1 1 3 8 1 1 1 1 ...
 $ v6 : int   2 7 2 3 2 7 2 2 2 2 ...
 $ v7 : chr   "1" "10" "2" "4" ...
 $ v8 : int   3 3 3 3 3 9 3 3 1 2 ...
 $ v9 : int   1 2 1 7 1 7 1 1 1 1 ...
 $ v10: int   1 1 1 1 1 1 1 1 5 1 ...
 $ v11: int   2 2 2 2 2 4 2 2 2 2 ...
> head(bcancer)
      v1 v2 v3 v4 v5 v6 v7 v8 v9 v10 v11
1 1000025  5  1  1  1  2  1  3  1  1  2
2 1002945  5  4  4  5  7 10  3  2  1  2
3 1015425  3  1  1  1  2  2  3  1  1  2
4 1016277  6  8  8  1  3  4  3  7  1  2
5 1017023  4  1  1  3  2  1  3  1  1  2
6 1017122  8 10 10  8  7 10  9  7  1  4
> tail(bcancer)
      v1 v2 v3 v4 v5 v6 v7 v8 v9 v10 v11
694 763235  3  1  1  1  2  1  2  1  2  2
695 776715  3  1  1  1  3  2  1  1  1  2
696 841769  2  1  1  1  2  1  1  1  1  2
697 888820  5 10 10  3  7  3  8 10  2  4
698 897471  4  8  6  4  3  4 10  6  1  4
699 897471  4  8  8  5  4  5 10  4  1  4
```

First, I dropped the first column (id) as it uniquely identifies each observation and would prevent the model from actually learning as opposed to memorizing the observations. Then, I added the column names for convenience purposes.

```
> bcancer <- bcancer[,-1] #Drop the ID column
> colnames(bcancer) <- c("thickness", "sizeunif", "shapeunif", "adhesion", "cellsize", "barenucl", "chromatin", "normnucleoli", "mitoses", "class")
```

Summary statistic for all variables did not show any irregularities:

```
> #summary stats for all variables
> summary(bcancer)
  thickness      sizeunif      shapeunif      adhesion      cellsize      barenucl1
Min.   : 1.000   Min.   : 1.000   Min.   : 1.000   Min.   : 1.000   Min.   : 1.000   Length:699
1st Qu.: 2.000   1st Qu.: 1.000   1st Qu.: 1.000   1st Qu.: 1.000   1st Qu.: 2.000   Class :character
Median : 4.000   Median : 1.000   Median : 1.000   Median : 1.000   Median : 2.000   Mode  :character
Mean   : 4.418   Mean   : 3.134   Mean   : 3.207   Mean   : 2.807   Mean   : 3.216
3rd Qu.: 6.000   3rd Qu.: 5.000   3rd Qu.: 5.000   3rd Qu.: 4.000   3rd Qu.: 4.000
Max.   :10.000   Max.   :10.000   Max.   :10.000   Max.   :10.000   Max.   :10.000
 chromatin  normnucleoli  mitoses      class
Min.   : 1.000   Min.   : 1.000   Min.   : 1.000   Min.   :2.00
1st Qu.: 2.000   1st Qu.: 1.000   1st Qu.: 1.000   1st Qu.:2.00
Median : 3.000   Median : 1.000   Median : 1.000   Median :2.00
Mean   : 3.438   Mean   : 2.867   Mean   : 1.589   Mean   :2.69
3rd Qu.: 5.000   3rd Qu.: 4.000   3rd Qu.: 1.000   3rd Qu.:4.00
Max.   :10.000   Max.   :10.000   Max.   :10.000   Max.   :4.00
```

Next, I converted the dependent variable (class) to factor and the bareneucle field to numeric. The last step introduced 16 NA values.

```
> bcancer$barenucl1 <- as.numeric(bcancer$barenucl1)
> bcancer$class <- as.factor(bcancer$class)
> sum(is.na(bcancer))
[1] 16
```

I also looked at the visual presentation of the data set using thickness and cellsize variables as an example:

```
> ###Visualize the data
>
> #using clump thickness and uniform cell size as independent variables
> qplot(data=bcancer, x=thickness, y=cellsize, color=class) + geom_jitter()
```

The graph below shows that while majority of the observation belonging to the class 2 (no cancer) are clustered together, it is not possible to clearly separate those two classes in this two-dimensional space.

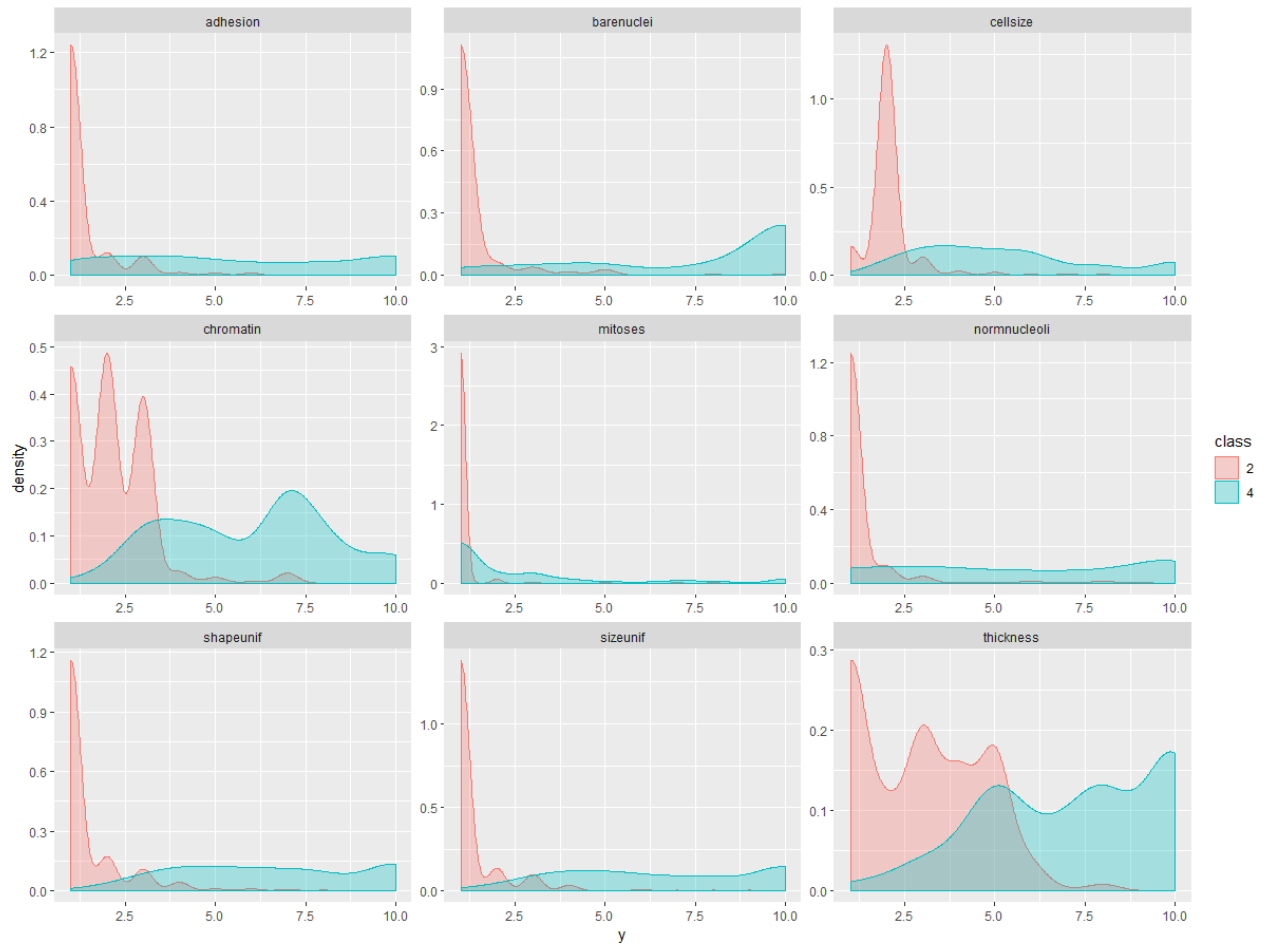


```
> #class=2 - benign; class = 4 cancer
> table(bcancer$class)
```

```
 2    4
458 241
```

There are total 458 benign observations in the data set and 241 observation with the cancer diagnosis and the density distribution of all independent variables shows that none of them can uniquely identify malignant observations.

```
> #density distribution for all independent variables by class
> gather(bcancer, x, y, thickness:mitoses) %>% ggplot(aes(x = y, color = class, fill = class)) +
+   geom_density(alpha = 0.3) + facet_wrap(~ x, scales = "free", ncol = 3)
```



In the next step, I imputed missing values in the bareuclei column with median values,

```
#make a copy of the df
> bcancer2 <- bcancer
>
> #impute missing values with the median value
> bcancer2$bareuclei <- ifelse(is.na(bcancer2$bareuclei), median(bcancer2$
bareuclei, na.rm = TRUE), bcancer2$bareuclei)
```

Finally, I split the data set into the training set (70%) and the testing set (30%).

```
> #split into training and testing data
> set.seed(123)
> trainIndex <- createDataPartition(bcancer2$class, p= 0.7, list = FALSE)
> bcancer_train <-bcancer2[trainIndex,]
> bcancer_test <- bcancer2[-trainIndex,]
```

As a result, the training part of the data set contains 490 observations and the testing part 209 observations.

I also checked proportions between the classes in the testing and training sets and compare them to the initial full dataset:

```
> #check class proportions
> prop.table(table(bcancer$class))

      2      4
0.6552217 0.3447783
> prop.table(table(bcancer_test$class))

      2      4
0.6555024 0.3444976
> prop.table(table(bcancer_train$class))

      2      4
0.655102 0.344898
```

In all three datasets, healthy observations (class2) represent about 65.5% and about 34.5% belongs to the observations with the breast cancer diagnosis.

In order to be able to use deep learning, I converted both training and testing data frames into h2o objects using the `as.h2o()` command:

```
> #####h2o#####
> ###Convert training and testing data frames into H2O objects
> cancer_h2otrain <- as.h2o(bcancer_train, key="train")
|=====| 100%
> cancer_h2otest <- as.h2o(bcancer_test, key="test")
|=====| 100%
>
> #check results
> class(cancer_h2otrain)
[1] "H2OFrame"
> class(cancer_h2otest)
[1] "H2OFrame"
```

These H2OFrames can be used to train and test deep learning models.

First, I used `h2o.deep learning()` command to train a model with three hidden layers with 50 nodes each using a hyperbolic tangent (tanh) activation function without regularization (dropout).

```
####dlmodel1 with 3 hidden layers 50 nodes without regularization (dropout)
> #train
> dlmodel1 <- h2o.deeplearning(x=1:9, y=10, training_frame = cancer_h2otrain, activation="Tanh",
hidden=c(50,50,50), epochs = 500)
|=====| 100%
```

As the output below shows, the H2O package provides a detailed list of model features and its performance on training data:

```
> dlmodel1
Model Details:
=====
```

```
H2OBinomialModel: deeplearning
Model ID: DeepLearning_model_R_1544812112618_242
Status of Neuron Layers: predicting class, 2-class classification, bernoulli
distribution, CrossEntropy loss, 5,702 weights/biases, 73.3 KB, 191,100 train
ing samples, mini-batch size 1
  layer units   type dropout    l1      l2 mean_rate rate_rms momentum m
ean_weight weight_rms mean_bias bias_rms
1      1      9 Input  0.00 %      NA      NA      NA      NA      NA
NA      NA      NA      NA      NA
2      2     50 Tanh  0.00 % 0.000000 0.000000  1.002337 0.004493 0.000000
-0.012986  0.223489 -0.007973 0.092561
3      3     50 Tanh  0.00 % 0.000000 0.000000  1.004175 0.000093 0.000000
0.004183  0.189595 -0.006429 0.135725
4      4     50 Tanh  0.00 % 0.000000 0.000000  1.004215 0.000012 0.000000
-0.002369  0.186843  0.008276 0.089226
5      5      2 Softmax      NA 0.000000 0.000000  1.004207 0.000015 0.000000
-0.077240  0.902894 -0.000555 0.117135
```

```
H2OBinomialMetrics: deeplearning
** Reported on training data. **
** Metrics reported on full training frame **
```

```
MSE: 1.025803e-08
RMSE: 0.0001012819
LogLoss: 6.729239e-06
Mean Per-Class Error: 0
AUC: 1
Gini: 1
```

Confusion Matrix (vertical: actual; across: predicted) for F1-optimal threshold:

	2	4	Error	Rate
2	321	0	0.000000	=0/321
4	0	169	0.000000	=0/169
Totals	321	169	0.000000	=0/490

Maximum Metrics: Maximum metrics at their respective thresholds

	metric	threshold	value	idx
1	max f1	0.999865	1.000000	21
2	max f2	0.999865	1.000000	21
3	max f0point5	0.999865	1.000000	21
4	max accuracy	0.999865	1.000000	21
5	max precision	1.000000	1.000000	0
6	max recall	0.999865	1.000000	21
7	max specificity	1.000000	1.000000	0

```

8           max absolute_mcc  0.999865 1.000000 21
9   max min_per_class_accuracy 0.999865 1.000000 21
10  max mean_per_class_accuracy 0.999865 1.000000 21

```

Gains/Lift Table: Extract with ``h2o.gainsLift(<model>, <data>)`` or ``h2o.gainsLift(<model>, valid=<T/F>, xval=<T/F>)``

On training data, the model showed 100% accuracy, which changed a little once I evaluated the model performance on the testing data set:

```

> #model performance on test data
> dlmodel1_perform <- h2o.performance(dlmodel1, cancer_h2otest)
> dlmodel1_perform
H2OBinomialMetrics: deeplearning

MSE: 0.05521148
RMSE: 0.2349712
LogLoss: 0.7167503
Mean Per-Class Error: 0.03649635
AUC: 0.9854015
Gini: 0.9708029

```

Confusion Matrix (vertical: actual; across: predicted) for F1-optimal threshold:

	2	4	Error	Rate
2	127	10	0.072993	=10/137
4	0	72	0.000000	=0/72
Totals	127	82	0.047847	=10/209

Maximum Metrics: Maximum metrics at their respective thresholds

		metric	threshold	value	idx
1		max f1	0.000002	0.935065	65
2		max f2	0.000002	0.972973	65
3		max f0point5	0.999930	0.915698	51
4		max accuracy	0.000002	0.952153	65
5		max precision	1.000000	1.000000	0
6		max recall	0.000002	1.000000	65
7		max specificity	1.000000	1.000000	0
8		max absolute_mcc	0.000002	0.902196	65
9	max min_per_class_accuracy		0.698953	0.941606	59
10	max mean_per_class_accuracy		0.000002	0.963504	65

Gains/Lift Table: Extract with ``h2o.gainsLift(<model>, <data>)`` or ``h2o.gainsLift(<model>, valid=<T/F>, xval=<T/F>)``

The overall accuracy showed 95.22% as only 10 observations out of 209 were misclassified. Moreover, the model mistakenly classifies 10 healthy observations as malignant, resulting in about 7.3% error for the class 2 and 100% accuracy for the class 4.

```
> h2o.confusionMatrix(dlmodel1_perform)
```



Confusion Matrix (vertical: actual; across: predicted) for max f1 @ threshold = 2.09912728463825e-06:

	2	4	Error	Rate
2	127	10	0.072993	=10/137
4	0	72	0.000000	=0/72
Totals	127	82	0.047847	=10/209

Due to the subject area, it is more important to minimize false negatives meaning to correctly identify all cases of malignant tumors so that the patients can receive an early diagnosis and treatment. If we consider cancer class as the target positive class in the sense of positive diagnosis of the disease, then situations, when a healthy patient is incorrectly referred for additional testing (false positives), are not ideal but are less dangerous. So, when choosing between the models, we cannot rely exclusively on overall accuracy as there are twice as many healthy samples as malignant in the training and testing data. It is necessary, in addition to overall accuracy, to prioritize the sensitivity, or true positive rate, or probability of correctly identifying all positive cancer diagnosis. Sensitivity for class 4, in this case, was 100% as all 72 cancerous observations were correctly identified.

Another useful indicator is the AUC (area under the curve) as it plots sensitivity (true positive rate) against the false positive rate.

For model evaluation and comparison purposes in the h2o package it is convenient to use an H2OModelMetrics object and its several accessor functions allowing to directly access particular indicators.

```
> h2o.mse(dlmodel1_perform)
[1] 0.05521148
> h2o.auc(dlmodel1_perform)
[1] 0.9854015
```

I followed the same procedures for a model with 3 hidden layers with 30 nodes each without regularization with the following output:

```

> ###dlmodel2 with 3 hidden layers 30 nodes each without regularization (dropout)
> #train
> dlmodel2 <- h2o.deeplearning(x=1:9, y=10, training_frame = cancer_h2otrain, activation="Tanh",
hidden=c(30,30,30), epochs = 500)
|=====
=====| 100%
> #model characteristics and performance on traing data
> dlmodel2
Model Details:
=====

H2OBinomialModel: deeplearning
Model ID: DeepLearning_model_R_1544812112618_246
Status of Neuron Layers: predicting class, 2-class classification, bernoulli distribution, CrossEntropy loss, 2,222 weights/biases, 31.8 KB, 245,000 training samples, mini-batch size 1
  layer units   type dropout    l1    l2 mean_rate rate_rms momentum mean_weight weight_rm
s mean_bias bias_rms
1     1     9 Input  0.00 %    NA    NA      NA      NA      NA      NA      NA      N
A      NA      NA
2     2    30 Tanh  0.00 % 0.000000 0.000000  1.004192 0.000061 0.000000   0.013869  0.28175
9 0.036869 0.123293
3     3    30 Tanh  0.00 % 0.000000 0.000000  1.004220 0.000006 0.000000  -0.009967  0.22349
8 0.020493 0.147806
4     4    30 Tanh  0.00 % 0.000000 0.000000  1.004223 0.000001 0.000000  -0.041165  0.24420
6 0.023283 0.136320
5     5     2 Softmax    NA 0.000000 0.000000  1.004224 0.000000 0.000000   0.011778  1.23446
7 -0.004679 0.191768

H2OBinomialMetrics: deeplearning
** Reported on training data. **
** Metrics reported on full training frame **

MSE:  3.964165e-09
RMSE: 6.296162e-05
LogLoss: 4.883045e-06
Mean Per-Class Error: 0
AUC: 1
Gini: 1

Confusion Matrix (vertical: actual; across: predicted) for F1-optimal threshold:
      2   4   Error   Rate
2    321   0 0.000000  =0/321
4      0 169 0.000000  =0/169
Totals 321 169 0.000000  =0/490

Maximum Metrics: Maximum metrics at their respective thresholds
      metric threshold   value idx
1      max f1  0.999485 1.000000  14
2      max f2  0.999485 1.000000  14
3      max f0point5 0.999485 1.000000  14
4      max accuracy 0.999485 1.000000  14
5      max precision 1.000000 1.000000   0
6      max recall  0.999485 1.000000  14
7      max specificity 1.000000 1.000000   0
8      max absolute_mcc 0.999485 1.000000  14
9      max min_per_class_accuracy 0.999485 1.000000  14

```

```
10 max mean_per_class_accuracy 0.999485 1.000000 14
```

Gains/Lift Table: Extract with ``h2o.gainsLift(<model>, <data>)`` or ``h2o.gainsLift(<model>, valid=<T/F>, xval=<T/F>)``

On the testing data the model demonstrated the following performance:

```
> #model performance on test data
> dlmodel2_perform <- h2o.performance(dlmodel2, cancer_h2otest)
> dlmodel2_perform
H2OBinomialMetrics: deeplearning
```

```
MSE: 0.06199205
RMSE: 0.248982
LogLoss: 0.9840077
Mean Per-Class Error: 0.03979116
AUC: 0.9832725
Gini: 0.966545
```

Confusion Matrix (vertical: actual; across: predicted) for F1-optimal threshold:

	2	4	Error	Rate
2	128	9	0.065693	=9/137
4	1	71	0.013889	=1/72
Totals	129	80	0.047847	=10/209

Maximum Metrics: Maximum metrics at their respective thresholds

	metric	threshold	value	idx
1	max f1	0.000121	0.934211	59
2	max f2	0.000000	0.970350	62
3	max f0point5	1.000000	0.922619	45
4	max accuracy	0.000121	0.952153	59
5	max precision	1.000000	1.000000	0
6	max recall	0.000000	1.000000	62
7	max specificity	1.000000	1.000000	0
8	max absolute_mcc	0.000121	0.899853	59
9	max min_per_class_accuracy	0.004123	0.934307	56
10	max mean_per_class_accuracy	0.000121	0.960209	59

Gains/Lift Table: Extract with ``h2o.gainsLift(<model>, <data>)`` or ``h2o.gainsLift(<model>, valid=<T/F>, xval=<T/F>)``

```
> ####evaluate
> #####h2o.ModelMetrics
>
> h2o.confusionMatrix(dlmodel2_perform)
Confusion Matrix (vertical: actual; across: predicted) for max f1 @ threshold = 1.6769586576830
9e-07:
      2  4  Error  Rate
2    126 11 0.080292 =11/137
4      1 71 0.013889 =1/72
Totals 127 82 0.057416 =12/209
>
> #h2o.performance() returns an H2OModelMetrics object
> #H2OModelMetrics object has several accessor functions
> h2o.mse(dlmodel2_perform)
[1] 0.05761159
> h2o.auc(dlmodel2_perform)
[1] 0.9821573
```

I continued to experiment with the model parameters and tested the following additional models (for brevity the output is presented in a following table):

- Model 3: 2 hidden layers 50 nodes each without regularization
- Model 4: 2 hidden layers 30 nodes each without regularization
- Model 5: 3 hidden layers with 50 nodes each with regularization to prevent overfitting
- Model 6: 3 hidden layers 30 nodes each with regularization to prevent overfitting
- Model 7: 2 hidden layers with 50 nodes each with regularization to prevent overfitting
- Model 8: 2 hidden layers 30 nodes each with regularization to prevent overfitting.

	<b>Model 1</b>	<b>Model 2</b>	<b>Model 3</b>	<b>Model 4</b>	<b>Model 5</b>	<b>Model6</b>	<b>Model 7</b>	<b>Model 8</b>
<b>Accuracy (%)</b>	95.22	94.26	94.73	94.74	97.61	97.61	97.61	97.13
<b>Sensitivity (%)</b>	100	98.61	97.22	100	97.22	95.83	98.61	100
<b>AUC</b>	0.9854015	0.9821573	0.9897607	0.9822739	0.9913828	0.9638078	0.9929035	0.9909773
<b>MSE</b>	0.05521148	0.05761159	0.05329071	0.06983468	0.037161	0.02815383	0.03272154	0.0298987

The above table shows that three models (models 5, 6, and 7) showed the best overall accuracy of 97.61%. However, they were not the best in detecting cancer diagnosis; models 1, 4, and 8 demonstrated 100% sensitivity. Model 7 demonstrated the best AUC (area under the curve) – 0.9929035, and model 6 had the smallest MSE. There was no one particular model which would show overall best results. Balancing the need for accuracy and prioritization of the correct cancer diagnosis, I would say that the model 8 was the best tool for detecting breast cancer.

In the absence of a clear answer, I was interested to see what model would be automatically recommended by the h2o package using random search, since it is not practical to search through all possible options. I used h2o functions `h2o.grid()` and `h2o.getGrid()` to automatically compare possible models given a set of parameters and the AUC choice criterion:.

```
> ####model tuning
> hyper_params <- list(
+   activation = c("Tanh", "TanhwithDropout"),
+   hidden = list(c(30,30), c(40, 40), c(50,50), c(30,30,30), c(40,40,40), c(50,50,50)),
+   input_dropout_ratio=c(0, 0.5),
+   rate = c(0.01, 0.25)
+ )
> search_criteria <- list(
+   strategy = "RandomDiscrete",
+   max_models=100, seed =123, stopping_rounds =5,
+   stopping_tolerance = 0.01
+ )
>
> randomSearch <- h2o.grid(
+   algorithm ="deeplearning",
+   grid_id="randomSearch",
+   training_frame = cancer_h2otrain,
+   validation_frame = cancer_h2otest,
+   x=1:9, y=10,
+   epochs=1,
+   stopping_metric ="misclassification",
+   hyper_params = hyper_params,
+   search_criteria = search_criteria
+ )
|=====| 100%
> search_criteria <- list(
+   strategy = "RandomDiscrete",
+   max_models=100, seed =123, stopping_rounds =5,
+   stopping_tolerance = 0.01
```

```

+ )
>
> randomSearch <- h2o.grid(
+   algorithm="deeplearning",
+   grid_id="randomSearch",
+   training_frame = cancer_h2otrain,
+   validation_frame = cancer_h2otest,
+   x=1:9, y=10,
+   epochs=1,
+   stopping_metric="misclassification",
+   hyper_params = hyper_params,
+   search_criteria = search_criteria
+ )
|=====
=====| 100%
> grid <- h2o.getGrid("randomSearch", sort_by="auc", decreasing = TRUE)
> grid
H2O Grid Details
=====
Grid ID: randomSearch
Used hyper parameters:
- activation
- hidden
- input_dropout_ratio
- rate
Number of models: 48
Number of failed models: 0

Hyper-Parameter Search Summary: ordered by decreasing auc
      activation      hidden input_dropout_ratio rate      model_ids      auc
1 TanhwithDropout [30, 30]          0.0 0.25 randomSearch_model_46 0.9946269261962692
2           Tanh   [30, 30]          0.0 0.01 randomSearch_model_18 0.994322789943228
3           Tanh   [30, 30]          0.5 0.01 randomSearch_model_2  0.9940186536901865
4           Tanh   [50, 50]          0.0 0.01 randomSearch_model_40 0.9932076236820764
5           Tanh [30, 30, 30]        0.0 0.25 randomSearch_model_12 0.9931062449310624

---
      activation      hidden input_dropout_ratio rate      model_ids      auc
43 TanhwithDropout [40, 40, 40]      0.5 0.25 randomSearch_model_45 0.9886455798864558
44           Tanh   [50, 50, 50]      0.5 0.01 randomSearch_model_29 0.9873276561232766
45           Tanh   [50, 50, 50]      0.0 0.01 randomSearch_model_38 0.987124898621249
46 TanhwithDropout [30, 30, 30]      0.5 0.01 randomSearch_model_31 0.9867193836171937
47 TanhwithDropout [40, 40, 40]      0.0 0.25 randomSearch_model_1  0.9863138686131386
48           Tanh   [50, 50]          0.5 0.01 randomSearch_model_44 0.9857055961070559

```

As the output above shows, the best suggested model is a relatively simpler one. It has two layers with 30 nodes each and a dropout rate of 0.25. I tested the suggested model on the test dataset and received the following results (partial output shown):

```

MSE:  0.02440406
RMSE: 0.156218
LogLoss: 0.08906241

```

Mean Per-Class Error: 0.01094891  
 AUC: 0.9946269  
 Gini: 0.9892539

Confusion Matrix (vertical: actual; across: predicted) for F1-optimal threshold:

	2	4	Error	Rate
2	134	3	0.021898	=3/137
4	0	72	0.000000	=0/72
Totals	134	75	0.014354	=3/209

It means that compared to the model that I manually chose in the previous step, the best model still demonstrated a 100% sensitivity to detecting cancer, but improved overall accuracy to 98.56%, while AUC increased to 0.9992539 from 0.9909773. It is actually a better model for the given data.